

[START FOR FREE](#)

CUSTOMER STORIES

TS IMAGINE ADOPTS GEN AI AT SCALE, SAVING 30% IN COSTS AND 4,000 HOURS OF EFFORT

With Snowflake's AI capabilities, TS Imagine unifies its data — as well as its data teams and technologies — across more than 500 clients, and implements generative AI at scale to improve efficiency and cut costs.


KEY RESULTS:

30%

Cost reduction by using Snowflake Cortex AI vs. other top external pretrained LLM APIs

4,000

Hours per year saved, previously spent on manual email monitoring tasks



[START FOR FREE](#)

TSIMAGINE

INDUSTRY

Financial Services

LOCATION

New York, London, Montreal, Hong Kong, Singapore,
Paris

SNOWFLAKE WORKLOADS USED



AI/ML



Applications



Collaboration



Data Engineering



Data Warehouse

CONTACT SALES

Solid data pays dividends

It's Monday morning and the financial markets are open. Across the globe, experts are logging on to their computers to start the work day. Portfolio managers. Heads of trading. Chief risk officers. The data they rely on to make decisions may impact the lives (and livelihoods) of millions of people. So that data better be right.

Those are the stakes TS Imagine handles on a daily basis for its customers. As a best-in-class SaaS platform for integrated electronic front-office trading, portfolio management and risk management, TS Imagine's offering comes with valuable out-of-the-box data for over 20 million financial instruments. Its data management team curates and normalizes that data so experts from over 500 institutions logging on every day — from multinational banks, huge hedge funds and more — can focus on generating returns within today's fast-evolving markets.



inefficiencies. Chief Operating and Data & Analytics Officer Thomas Bodenski knew he needed to automate, unify and simplify data and processes, in part to better equip the organization with generative AI (gen AI) tools.

Over the last three years, as TS Imagine executed its data strategy, Bodenski overcame these challenges by replacing the company's homegrown solutions with Snowflake. In turning to the AI Data Cloud and Cortex AI, his global Data & Analytics office has streamlined operations while successfully rolling out gen AI use cases across the enterprise, saving time and democratizing access to information.

Story Highlights

- **Easy-to-build AI use cases for better productivity:** Snowflake's platform is so easy to use that, in less than a week, a team member without an AI engineering background can roll out AI use cases that improve companywide workflows.
- **Improved efficiency for faster insights:** With generative AI at scale, business users can gather information far more quickly by conversing with an easy-to-use chatbot built on Streamlit in Snowflake.
- **Massive time savings for employees:** By implementing a retrieval augmented generation (RAG)-based process on Snowflake, near-instant, AI-powered categorization allows TS Imagine to easily save 4,000 annual hours previously spent on manual email monitoring tasks.

The real ROI of generative AI at scale

Getting buy-in on Snowflake was tricky at first. Some of Bodenski's colleagues were skeptical they could accomplish everything they wanted with one platform. With help from Snowflake, TS Imagine piloted a demo analyzing, normalizing and sharing aggregated data based on hundreds of millions of prices and price interactions. This was a level of data they hadn't been able to tackle without Snowflake. The TS Imagine team was sold.

Secure data sharing became a cinch with clients thanks to a private Snowflake Marketplace listing, and data volume was no longer a pain point. "Every week, we grow by almost one billion rows," says Bodenski. "With Snowflake, those types of volumes don't scare us. So what? Of course we can handle that."

But what has proven even more game-changing is TS Imagine's implementation of AI at scale. Bodenski and Global Head of AI Engineering Cathy Siu knew that generative AI would allow the team to turn all its unstructured data, like emails, into structured data that their machine learning systems could act on. In particular, retrieval augmented generation, which combines the powerful generative capabilities of large language models (LLMs) with a retrieval mechanism that sources up-to-date, contextually relevant information, could help their business. So Siu worked on a RAG use case.

Initially she used GPT-4-Turbo and Chroma DB; precision rates soared, but it still required a lot of prompt engineering, trial and error, and help from the engineering team to operationalize the use case. Once Cortex AI was available, she transferred the use case over in under a week. Snowflake's platform was easy enough to use



“Doing everything exclusively in Snowflake was game-changing. Now we design something on a Thursday, and by Tuesday it’s in production.”

Thomas Bodenski

COO & Chief Data & Analytics Officer, TS Imagine

With everything they need to build, manage and run gen AI or RAG-based use cases, the data analytics team now accomplishes more with the same spend and headcount. Employees are free to focus on strategic tasks as opposed to manual data prep and email sorting. And since there’s no need to move or share data outside Snowflake’s secure, controlled environment, TS Imagine maintains the strict data privacy and protection its clients expect — and the product quality they deserve.

Being bullish on efficiency, saving 4,000 hours a year

For TS Imagine, email absorption is mission-critical. The company receives an estimated 100,000 emails a year, which include critical notifications from data providers about upcoming changes to their data products. Receiving that information allows TS Imagine to test and prepare its products. “If we’re not ready for these changes, there’s going to be a production outage,” Bodenski says. Instead of producing insights, members of the data management team had to read those emails manually. Every single one. And even then, sometimes mistakes would happen.

For its first RAG-based use case, TS Imagine automated email intake using Cortex AI. Instead of spending 4,000 hours on an error-prone sorting process, the AI deletes duplicate or non-relevant emails, and creates, assigns, prioritizes and schedules JIRA tickets for them. The team hasn’t missed a notification since implementing the process in December 2023.

Investing in better customer experiences

[START FOR FREE](#)

effectively triage, respond to and resolve customer-facing issues much faster.

“With Snowflake, I can empower smart people to bring AI to life in one place. Cortex is a one-stop shop. It scales, it’s easy, and the data stays 100% in Snowflake’s environment.”

Thomas Bodenski

COO & Chief Data & Analytics Officer, TS Imagine

And rather than going through the painstaking process of reading PDFs detailing terms and conditions for securities available within its trading and risk management products, TS Imagine uses Snowflake’s large language model Arctic, along with other LLMs such as Mistral and Llama, to convert documents into structured data, automatically classifying, summarizing and parsing them.

Surfacing a wealth of knowledge

Though TS Imagine is only a few years old, it’s the product of a merger between two companies with three decades of documentation, processes and information for employees to sift through. Whether it was a tenured manager looking for straightforward answers or a new hire completing onboarding, having so many places to look for company data was tiresome.

To better surface information and fuel employee productivity, Siu ran the tens of thousands of documents, articles, support tickets and other items that comprise the company’s knowledge base through Snowflake Cortex AI, then built a chatbot with Streamlit in Snowflake. Now employees can ask questions in a handy, user-friendly interface, cutting down time previously wasted hunting for answers and bridging the institutional knowledge gap.

AI as the ultimate asset

With Snowflake, Bodenski and his team have already done what almost 90% of companies have yet to do, [according to McKinsey](#): Adopt gen AI at scale. But they’re just getting started. Given how straightforward it is to develop and deploy on Snowflake, creativity is their only limit.



START FOR FREE



4

Business days to operationalize new RAG use cases

RELATED CUSTOMER STORIES



tapestry

Tapestry Builds Its Own Analytics Platform on Snowflake

[Read the Customer Story >](#)



START FOR FREE



bond.

Bond Brand Loyalty Saves 50% in Costs by Migrating from Managed Spark to Snowflake

[Read the Customer Story >](#)



petco

Petco Unleashes New Revenue Opportunities and Savings

[Read the Customer Story >](#)

[START FOR FREE](#)

Try Snowflake free for 30 days and experience the AI Data Cloud that helps eliminate the complexity, cost and constraints inherent with other solutions.

[START FOR FREE](#)[VIEW PRICING](#)[LIVE DEMO](#)[Privacy Notice](#)[Site Terms](#)[Cookie Settings](#)[Do Not Share My Personal Information](#)

© 2025 Snowflake Inc. All Rights Reserved | If you'd rather not receive future emails from Snowflake, [unsubscribe here](#) or [customize your communication preferences](#)